

ict-employee-absenteeism-from-work

July 6, 2025

1 Employee Absenteeism

The database was created with records of absenteeism at work from July 2007 to July 2010 at a courier company in Brazil.

2 Aim

The exercise will address **Absenteeism** at a company during work time.

2.1 Problem

The problem is that the work environment of today is more: - Competitive - Managers set unachievable business goals - have an elevated risk of becoming unemployed

This can lead to an increase in pressure and stress of the employee. Those factors influence employee health, which is of course undesirable.

2.2 What is Absenteeism?

Absence from work during normal working hours resulting in temporary incapacity to execute a regular working activity. - Based on what information should we predict whether an employee is expected to be absent or not? - How should we measure absenteeism?

2.3 Purpose of the business exercise:

Explore whether a person presenting certain characteristics is expected to be away from work at some point in time or not.

We want to know for how many working hours any employee could be away from work based on information like: - How far they live from their workplace. - How many children and pets they have. - Do they have higher education?

```
[1]: !pip install -q hvplot
!pip install -q xlrd
```

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import hvplot.pandas
```

```
[2]: data = pd.read_excel('D:/Project4/Absent_at_work.xls')
data.columns = data.columns.str.replace(" ", "_")

data.head()
```

```
[2]:
```

	ID	Reason_for_absence	Month_of_absence	Day_of_the_week	Seasons	\
0	11	26.0	7.0	3	1	
1	36	0.0	7.0	3	1	
2	3	23.0	7.0	4	1	
3	7	7.0	7.0	5	1	
4	11	23.0	7.0	5	1	

	Transportation_expense	Distance_from_Residence_to_Work	Service_time	\
0	289.0	36.0	13.0	
1	118.0	13.0	18.0	
2	179.0	51.0	18.0	
3	279.0	5.0	14.0	
4	289.0	36.0	13.0	

	Age	Work_load_Average/day_	...	Disciplinary_failure	Education	Son	\
0	33.0	239554.0	...	0.0	1.0	2.0	
1	50.0	239554.0	...	1.0	1.0	1.0	
2	38.0	239554.0	...	0.0	1.0	0.0	
3	39.0	239554.0	...	0.0	1.0	2.0	
4	33.0	239554.0	...	0.0	1.0	2.0	

	Social_drinker	Social_smoker	Pet	Weight	Height	Body_mass_index	\
0	1.0	0.0	1.0	90.0	172.0	30.0	
1	1.0	0.0	0.0	98.0	178.0	31.0	
2	1.0	0.0	0.0	89.0	170.0	31.0	
3	1.0	1.0	0.0	68.0	168.0	24.0	
4	1.0	0.0	1.0	90.0	172.0	30.0	

	Absenteeism_time_in_hours
0	4.0
1	0.0
2	2.0
3	4.0

[5 rows x 21 columns]

3 Exploratory Data Analysis (EDA)

```
[3]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 740 entries, 0 to 739
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                    740 non-null    int64
1   Reason_for_absence                   737 non-null    float64
2   Month_of_absence                     739 non-null    float64
3   Day_of_the_week                      740 non-null    int64
4   Seasons                             740 non-null    int64
5   Transportation_expense               733 non-null    float64
6   Distance_from_Residence_to_Work     737 non-null    float64
7   Service_time                        737 non-null    float64
8   Age                                  737 non-null    float64
9   Work_load_Average/day_              730 non-null    float64
10  Hit_target                          734 non-null    float64
11  Disciplinary_failure                 734 non-null    float64
12  Education                           730 non-null    float64
13  Son                                  734 non-null    float64
14  Social_drinker                      737 non-null    float64
15  Social_smoker                       736 non-null    float64
16  Pet                                  738 non-null    float64
17  Weight                              739 non-null    float64
18  Height                              726 non-null    float64
19  Body_mass_index                     709 non-null    float64
20  Absenteeism_time_in_hours            718 non-null    float64
dtypes: float64(18), int64(3)
memory usage: 121.5 KB
```

```
[4]: data = data.dropna()
```

The data doesn't have any missing values.

```
[5]: pd.set_option("display.float_format", "{:.2f}".format)
data.describe()
```

```
[5]:
```

	ID	Reason_for_absence	Month_of_absence	Day_of_the_week	Seasons	\
count	639.00	639.00	639.00	639.00	639.00	

mean	17.73	19.20	6.16	3.89	2.52
std	11.04	8.49	3.34	1.43	1.09
min	1.00	0.00	0.00	2.00	1.00
25%	8.00	13.00	3.00	3.00	2.00
50%	18.00	23.00	6.00	4.00	3.00
75%	28.00	27.00	9.00	5.00	3.00
max	36.00	28.00	12.00	6.00	4.00

	Transportation_expense	Distance_from_Residence_to_Work	Service_time	\
count	639.00	639.00	639.00	
mean	221.09	29.67	12.73	
std	64.97	14.70	4.36	
min	118.00	5.00	1.00	
25%	179.00	16.00	9.00	
50%	225.00	26.00	13.00	
75%	260.00	50.00	16.00	
max	388.00	52.00	29.00	

	Age	Work_load_Average/day_	...	Disciplinary_failure	Education	\
count	639.00	639.00	...	639.00	639.00	
mean	36.69	270782.13	...	0.05	1.31	
std	6.57	39049.26	...	0.22	0.69	
min	27.00	205917.00	...	0.00	1.00	
25%	31.00	244387.00	...	0.00	1.00	
50%	37.00	264249.00	...	0.00	1.00	
75%	40.00	284853.00	...	0.00	1.00	
max	58.00	378884.00	...	1.00	4.00	

	Son	Social_drinker	Social_smoker	Pet	Weight	Height	\
count	639.00	639.00	639.00	639.00	639.00	639.00	
mean	1.02	0.57	0.07	0.74	79.31	172.14	
std	1.08	0.49	0.26	1.32	12.95	6.05	
min	0.00	0.00	0.00	0.00	56.00	163.00	
25%	0.00	0.00	0.00	0.00	69.00	169.00	
50%	1.00	1.00	0.00	0.00	83.00	170.00	
75%	2.00	1.00	0.00	1.00	89.00	172.00	
max	4.00	1.00	1.00	8.00	108.00	196.00	

	Body_mass_index	Absenteeism_time_in_hours
count	639.00	639.00
mean	26.77	7.02
std	4.34	13.81
min	19.00	0.00
25%	24.00	2.00
50%	25.00	3.00
75%	31.00	8.00
max	38.00	120.00

[8 rows x 21 columns]

```
[6]: for column in data.columns:
      print(f"=====Column: {column}=====")
      print(f"Number of unique values: {data[column].nunique()}")
      print(f"Max: {data[column].max()}")
      print(f"Min: {data[column].min()}")
```

```
=====Column: ID=====
Number of unique values: 36
Max: 36
Min: 1
=====Column: Reason_for_absence=====
Number of unique values: 27
Max: 28.0
Min: 0.0
=====Column: Month_of_absence=====
Number of unique values: 13
Max: 12.0
Min: 0.0
=====Column: Day_of_the_week=====
Number of unique values: 5
Max: 6
Min: 2
=====Column: Seasons=====
Number of unique values: 4
Max: 4
Min: 1
=====Column: Transportation_expense=====
Number of unique values: 24
Max: 388.0
Min: 118.0
=====Column: Distance_from_Residence_to_Work=====
Number of unique values: 25
Max: 52.0
Min: 5.0
=====Column: Service_time=====
Number of unique values: 18
Max: 29.0
Min: 1.0
=====Column: Age=====
Number of unique values: 22
Max: 58.0
Min: 27.0
=====Column: Work_load_Average/day_=====
Number of unique values: 38
```

```

Max: 378884.0
Min: 205917.0
=====Column: Hit_target=====
Number of unique values: 13
Max: 100.0
Min: 81.0
=====Column: Disciplinary_failure=====
Number of unique values: 2
Max: 1.0
Min: 0.0
=====Column: Education=====
Number of unique values: 4
Max: 4.0
Min: 1.0
=====Column: Son=====
Number of unique values: 5
Max: 4.0
Min: 0.0
=====Column: Social_drinker=====
Number of unique values: 2
Max: 1.0
Min: 0.0
=====Column: Social_smoker=====
Number of unique values: 2
Max: 1.0
Min: 0.0
=====Column: Pet=====
Number of unique values: 6
Max: 8.0
Min: 0.0
=====Column: Weight=====
Number of unique values: 26
Max: 108.0
Min: 56.0
=====Column: Height=====
Number of unique values: 14
Max: 196.0
Min: 163.0
=====Column: Body_mass_index=====
Number of unique values: 17
Max: 38.0
Min: 19.0
=====Column: Absenteeism_time_in_hours=====
Number of unique values: 19
Max: 120.0
Min: 0.0

```

```
[7]: data.columns
```

```
[7]: Index(['ID', 'Reason_for_absence', 'Month_of_absence', 'Day_of_the_week',  
         'Seasons', 'Transportation_expense', 'Distance_from_Residence_to_Work',  
         'Service_time', 'Age', 'Work_load_Average/day_', 'Hit_target',  
         'Disciplinary_failure', 'Education', 'Son', 'Social_drinker',  
         'Social_smoker', 'Pet', 'Weight', 'Height', 'Body_mass_index',  
         'Absenteeism_time_in_hours'],  
        dtype='object')
```

```
[8]: data.ID.nunique()
```

```
[8]: 36
```

```
[9]: data.ID.value_counts().hvplot.hist(bins=data.ID.nunique(), height=300,  
    ↪width=500)
```

```
[9]: :Histogram    [count]    (Count)
```

ID: individual identification (in this case we have 34 employees) indicates precisely who has been away during working hours. Will this information improve our analysis in any way? No, because it's only a label variable (a number that is there to distinguish the individuals from one another, not to carry any numeric information).

So we are going to drop this column

```
[10]: data.drop('ID', axis=1, inplace=True)
```

```
[13]: for col in cat_col:  
      print(col)
```

```
Month_of_absence  
Day_of_the_week  
Seasons  
Service_time  
Hit_target  
Disciplinary_failure  
Education  
Son  
Social_drinker  
Social_smoker  
Pet  
Height  
Body_mass_index  
Absenteeism_time_in_hours
```

Reason for Absence: We have 28 reason of absence from 0 to 28.

```
[12]: cat_col = [col for col in data.columns if data[col].nunique() < 20]

      for col in cat_col:
          globals()[col] = data[col].value_counts().hvplot.bar(height=350, width=350)
```

```
[14]: Education
```

```
[14]: :Bars    [Education]    (count)
```

```
[15]: cat_col[0]
```

```
[15]: 'Month_of_absence'
```

```
[16]: Month_of_absence
```

```
[16]: :Bars    [Month_of_absence]    (count)
```

```
[17]: day_list=data['Day_of_the_week'].tolist()
      from functools import reduce

      total_days = reduce((lambda x, y: x + y), day_list)
      total_days
```

```
[17]: 2484
```

```
[18]: # Visualizing the distribution of the data for every feature
      data.hvplot.hist(y='Reason_for_absence', height=350, width=350)
```

```
[18]: :Histogram    [Reason_for_absence]    (Count)
```

```
[19]: Month_of_absence
```

```
[19]: :Bars    [Month_of_absence]    (count)
```

```
[20]: data.hvplot.hist(y='Month_of_absence', bins=12, width=350, height=350)
```

```
[20]: :Histogram    [Month_of_absence]    (Count)
```

```
[21]: data['Day_of_the_week'].value_counts().hvplot.bar(height=350, width=350)
```

```
[21]: :Bars    [Day_of_the_week]    (count)
```

```
[22]: data['Seasons'].value_counts().hvplot.bar(height=350, width=350)
```

```
[22]: :Bars    [Seasons]    (count)
```

```
[23]: data.columns
```



```
[23]: Index(['Reason_for_absence', 'Month_of_absence', 'Day_of_the_week', 'Seasons',
        'Transportation_expense', 'Distance_from_Residence_to_Work',
        'Service_time', 'Age', 'Work_load_Average/day_', 'Hit_target',
        'Disciplinary_failure', 'Education', 'Son', 'Social_drinker',
        'Social_smoker', 'Pet', 'Weight', 'Height', 'Body_mass_index',
        'Absenteeism_time_in_hours'],
        dtype='object')
```

```
[24]: print(f"{data['Absenteeism_time_in_hours'].value_counts()}")
```

```
Absenteeism_time_in_hours
8.00      178
2.00      136
3.00       99
1.00       78
4.00       52
0.00       34
16.00      17
24.00      15
40.00       7
5.00        6
80.00        3
120.00       3
32.00        3
112.00       2
64.00        2
7.00         1
56.00        1
104.00       1
48.00        1
Name: count, dtype: int64
```

```
[25]: print(data['Absenteeism_time_in_hours'].value_counts()[0])
```

```
34
```

```
[26]: print(data['Absenteeism_time_in_hours'].value_counts()[1])
```

```
78
```

```
[27]: print(f"{data['Absenteeism_time_in_hours'].value_counts().iloc[0]/\n
↪data['Absenteeism_time_in_hours'].value_counts().iloc[1]}")
```

```
1.3088235294117647
```

```
[28]: data["Reason_for_absence"].value_counts()
```

```
[28]: Reason_for_absence
      23.00    119
      28.00    102
      27.00     60
      13.00     51
       0.00     34
      19.00     33
      26.00     30
      25.00     30
      22.00     29
      11.00     24
      10.00     22
      18.00     16
      14.00     16
       1.00     16
       7.00     15
      12.00      8
       6.00      6
      21.00      5
       8.00      5
      24.00      3
       5.00      3
       9.00      3
      16.00      3
      15.00      2
       4.00      2
       3.00      1
       2.00      1
      Name: count, dtype: int64
```

```
[29]: data["Reason_for_absence"] = data["Reason_for_absence"].map(
      {
      0: "Group_1", 1: "Group_1", 2: "Group_1", 3: "Group_1",
      4: "Group_1", 5: "Group_1", 6: "Group_1", 7: "Group_1",
      8: "Group_1", 9: "Group_1", 10: "Group_1", 11: "Group_1",
      12: "Group_1", 13: "Group_1", 14: "Group_1", 15: "Group_2",
      16: "Group_2", 17: "Group_2", 17: "Group_2", 18: "Group_3",
      19: "Group_3", 20: "Group_3", 21: "Group_3", 22: "Group_4",
      23: "Group_4", 24: "Group_4", 25: "Group_4", 26: "Group_4",
      27: "Group_4", 28: "Group_4"
      }
      )
      # data["Reason for Absence"] = data["Reason for Absence"].astype("category").
      ↪ cat.codes
      data["Reason_for_absence"].value_counts()
```

```
[29]: Reason_for_absence
      Group_4      373
      Group_1      207
      Group_3       54
      Group_2        5
      Name: count, dtype: int64
```

```
[30]: data_1 = pd.get_dummies(data, columns=['Reason_for_absence'])
```

```
[31]: data_1.head()
```

```
[31]:   Month_of_absence  Day_of_the_week  Seasons  Transportation_expense \
0                7.00                3         1                289.00
1                7.00                3         1                118.00
2                7.00                4         1                179.00
3                7.00                5         1                279.00
4                7.00                5         1                289.00
```

```
      Distance_from_Residence_to_Work  Service_time  Age \
0                36.00                13.00  33.00
1                13.00                18.00  50.00
2                51.00                18.00  38.00
3                 5.00                14.00  39.00
4                36.00                13.00  33.00
```

```
      Work_load_Average/day_  Hit_target  Disciplinary_failure  ... \
0                239554.00        97.00                0.00  ...
1                239554.00        97.00                1.00  ...
2                239554.00        97.00                0.00  ...
3                239554.00        97.00                0.00  ...
4                239554.00        97.00                0.00  ...
```

```
      Social_smoker  Pet  Weight  Height  Body_mass_index \
0                0.00  1.00   90.00  172.00                30.00
1                0.00  0.00   98.00  178.00                31.00
2                0.00  0.00   89.00  170.00                31.00
3                1.00  0.00   68.00  168.00                24.00
4                0.00  1.00   90.00  172.00                30.00
```

```
      Absenteeism_time_in_hours  Reason_for_absence_Group_1 \
0                4.00                False
1                0.00                True
2                2.00                False
3                4.00                True
4                2.00                False
```

```
      Reason_for_absence_Group_2  Reason_for_absence_Group_3 \
```

0	False	False
1	False	False
2	False	False
3	False	False
4	False	False

Reason_for_absence_Group_4	
0	True
1	False
2	True
3	False
4	True

[5 rows x 23 columns]

```
[32]: data_1.dtypes
```

```
[32]: Month_of_absence      float64
      Day_of_the_week       int64
      Seasons               int64
      Transportation_expense float64
      Distance_from_Residence_to_Work float64
      Service_time          float64
      Age                   float64
      Work_load_Average/day_ float64
      Hit_target            float64
      Disciplinary_failure   float64
      Education             float64
      Son                   float64
      Social_drinker        float64
      Social_smoker         float64
      Pet                   float64
      Weight                float64
      Height                float64
      Body_mass_index       float64
      Absenteeism_time_in_hours float64
      Reason_for_absence_Group_1      bool
      Reason_for_absence_Group_2      bool
      Reason_for_absence_Group_3      bool
      Reason_for_absence_Group_4      bool
      dtype: object
```

```
[33]: data_1.dropna(inplace=True)
      data_1.isna().sum()
```

```
[33]: Month_of_absence      0
      Day_of_the_week       0
```

Seasons	0
Transportation_expense	0
Distance_from_Residence_to_Work	0
Service_time	0
Age	0
Work_load_Average/day_	0
Hit_target	0
Disciplinary_failure	0
Education	0
Son	0
Social_drinker	0
Social_smoker	0
Pet	0
Weight	0
Height	0
Body_mass_index	0
Absenteeism_time_in_hours	0
Reason_for_absence_Group_1	0
Reason_for_absence_Group_2	0
Reason_for_absence_Group_3	0
Reason_for_absence_Group_4	0

dtype: int64

```
[34]: data_1["Education"] = data_1.Education.map({1: 0, 2: 1, 3: 1, 4: 1})
```

```
[35]: data_1.Education.value_counts()
```

```
[35]: Education
0      522
1      117
Name: count, dtype: int64
```

```
[36]: print(f"Null Values Sum: {data_1.Education.isna().sum()}")
```

Null Values Sum: 0

```
[37]: data_2 = pd.get_dummies(data_1, columns=["Education"], drop_first=True)
data_2.columns
```

```
[37]: Index(['Month_of_absence', 'Day_of_the_week', 'Seasons',
        'Transportation_expense', 'Distance_from_Residence_to_Work',
        'Service_time', 'Age', 'Work_load_Average/day_', 'Hit_target',
        'Disciplinary_failure', 'Son', 'Social_drinker', 'Social_smoker', 'Pet',
        'Weight', 'Height', 'Body_mass_index', 'Absenteeism_time_in_hours',
        'Reason_for_absence_Group_1', 'Reason_for_absence_Group_2',
        'Reason_for_absence_Group_3', 'Reason_for_absence_Group_4',
        'Education_1'],
```

```
dtype='object')
```

```
[38]: data_2.head()
```

```
[38]:   Month_of_absence  Day_of_the_week  Seasons  Transportation_expense  \
0                7.00                3         1                289.00
1                7.00                3         1                118.00
2                7.00                4         1                179.00
3                7.00                5         1                279.00
4                7.00                5         1                289.00

      Distance_from_Residence_to_Work  Service_time  Age  \
0                36.00                13.00  33.00
1                13.00                18.00  50.00
2                51.00                18.00  38.00
3                 5.00                14.00  39.00
4                36.00                13.00  33.00

      Work_load_Average/day_  Hit_target  Disciplinary_failure  ...  Pet  Weight  \
0                239554.00        97.00                0.00  ...  1.00   90.00
1                239554.00        97.00                1.00  ...  0.00   98.00
2                239554.00        97.00                0.00  ...  0.00   89.00
3                239554.00        97.00                0.00  ...  0.00   68.00
4                239554.00        97.00                0.00  ...  1.00   90.00

      Height  Body_mass_index  Absenteeism_time_in_hours  \
0    172.00                30.00                4.00
1    178.00                31.00                0.00
2    170.00                31.00                2.00
3    168.00                24.00                4.00
4    172.00                30.00                2.00


      Reason_for_absence_Group_1  Reason_for_absence_Group_2  \
0                False                False
1                True                False
2                False                False
3                True                False
4                False                False

      Reason_for_absence_Group_3  Reason_for_absence_Group_4  Education_1
0                False                True                False
1                False                False                False
2                False                True                False
3                False                False                False
4                False                True                False
```

```
[5 rows x 23 columns]
```


```
[39]: plt.figure(figsize=(20, 15))
sns.heatmap(data_2.corr(), annot=True)
```

[39]: <Axes: >



```
[40]: data_2['Absenteeism_time_in_hours'].
hist(bins=data_2['Absenteeism_time_in_hours'].nunique())
```

[40]: <Axes: >



4 Applying machine learning algorithms

```
[41]: import pickle

from sklearn.model_selection import train_test_split, cross_val_score, \
    GridSearchCV
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.pipeline import Pipeline

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, accuracy_score, \
    classification_report
from sklearn.metrics import precision_recall_curve, roc_curve, roc_auc_score
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
from sklearn.svm import SVC

X = data_2.drop('Absenteeism_time_in_hours', axis=1)
y = np.where(data_2["Absenteeism_time_in_hours"] > \
    data_2["Absenteeism_time_in_hours"].median(), 1, 0)
```



```

print(f'X shape: {X.shape}')
print(f'y shape: {y.shape}')

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=42)

pipe = Pipeline([
    ('min_max_scaler', MinMaxScaler()),
    ('std_scaler', StandardScaler())
])

X_train = pipe.fit_transform(X_train)
X_test = pipe.transform(X_test)

```

```

X shape: (639, 22)
y shape: (639,)

```

```

[42]: def evaluate(model, X_train, X_test, y_train, y_test):
    y_test_pred = model.predict(X_test)
    y_train_pred = model.predict(X_train)

    print("TRAINIG RESULTS: \n=====")
    clf_report = pd.DataFrame(classification_report(y_train, y_train_pred,
↳output_dict=True))
    print(f"CONFUSION MATRIX:\n{confusion_matrix(y_train, y_train_pred)}")
    print(f"ACCURACY SCORE:\n{accuracy_score(y_train, y_train_pred):.4f}")
    print(f"CLASSIFICATION REPORT:\n{clf_report}")

    print("TESTING RESULTS: \n=====")
    clf_report = pd.DataFrame(classification_report(y_test, y_test_pred,
↳output_dict=True))
    print(f"CONFUSION MATRIX:\n{confusion_matrix(y_test, y_test_pred)}")
    print(f"ACCURACY SCORE:\n{accuracy_score(y_test, y_test_pred):.4f}")
    print(f"CLASSIFICATION REPORT:\n{clf_report}")

```

5 Logistic Regression

```

[43]: print("\n=====LOGISTIC REGRESSION=====")
lr_clf = LogisticRegression(solver='liblinear', penalty='l1')
lr_clf.fit(X_train, y_train)
evaluate(lr_clf, X_train, X_test, y_train, y_test)

```

```

=====LOGISTIC REGRESSION=====
TRAINIG RESULTS:

```

```
=====
CONFUSION MATRIX:
[[199  49]
 [ 54 145]]
ACCURACY SCORE:
0.7696
CLASSIFICATION REPORT:
              0          1  accuracy  macro avg  weighted avg
precision    0.79    0.75      0.77      0.77      0.77
recall       0.80    0.73      0.77      0.77      0.77
f1-score     0.79    0.74      0.77      0.77      0.77
support     248.00  199.00      0.77    447.00    447.00
TESTING RESULTS:
=====
```

```
CONFUSION MATRIX:
[[84 15]
 [36 57]]
ACCURACY SCORE:
0.7344
CLASSIFICATION REPORT:
              0          1  accuracy  macro avg  weighted avg
precision    0.70    0.79      0.73      0.75      0.74
recall       0.85    0.61      0.73      0.73      0.73
f1-score     0.77    0.69      0.73      0.73      0.73
support      99.00  93.00      0.73    192.00    192.00
```

```
[44]: def plot_precision_recall_vs_threshold(precisions, recalls, thresholds):
    plt.plot(thresholds, precisions[:-1], "b--", label="Precision")
    plt.plot(thresholds, recalls[:-1], "g--", label="Recall")
    plt.xlabel("Threshold")
    plt.legend(loc="upper left")
    plt.title("Precision/Recall Tradeoff")

def plot_roc_curve(fpr, tpr, label=None):
    plt.plot(fpr, tpr, linewidth=2, label=label)
    plt.plot([0, 1], [0, 1], "k--")
    plt.axis([0, 1, 0, 1])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve')

precisions, recalls, thresholds = precision_recall_curve(y_test, lr_clf.
    ↪predict(X_test))
plt.figure(figsize=(14, 25))
plt.subplot(4, 2, 1)
```

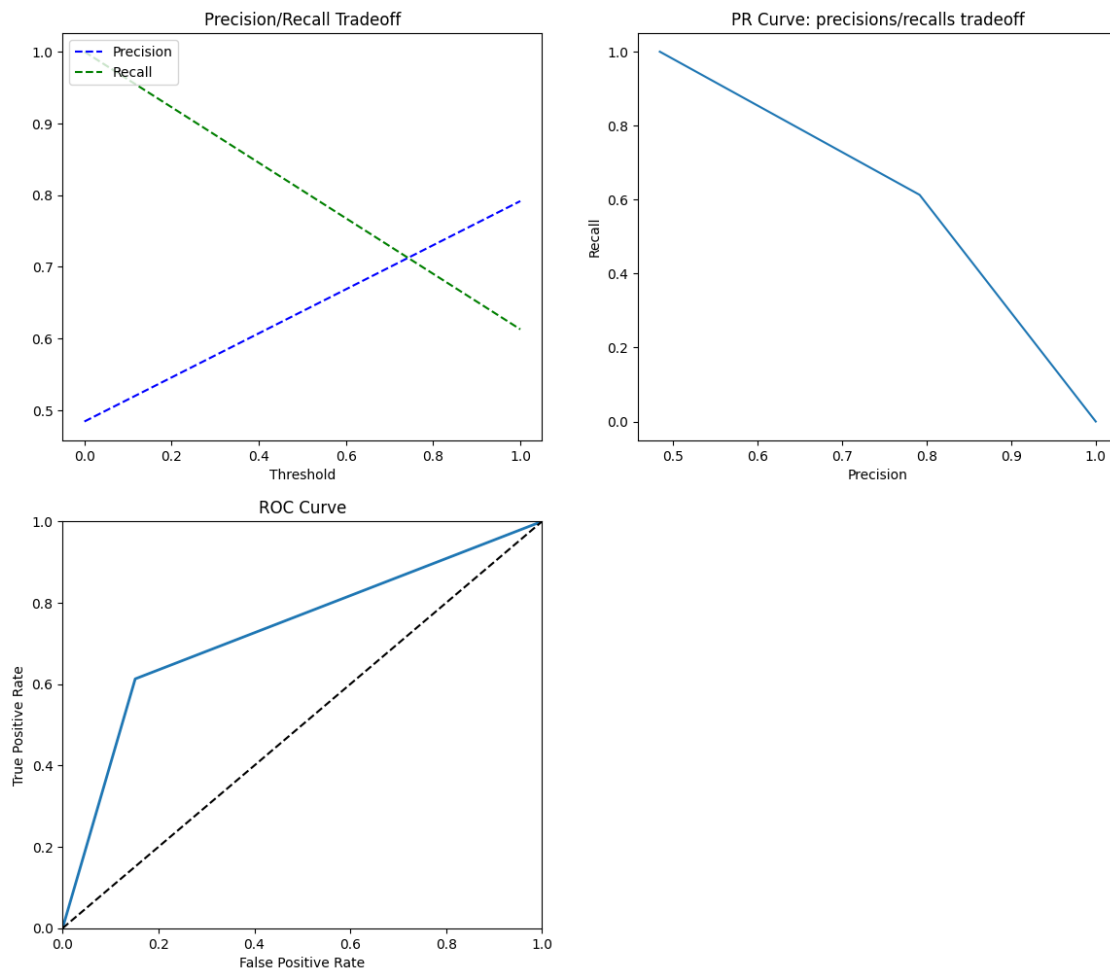
```

plot_precision_recall_vs_threshold(precisions, recalls, thresholds)

plt.subplot(4, 2, 2)
plt.plot(precisions, recalls)
plt.xlabel("Precision")
plt.ylabel("Recall")
plt.title("PR Curve: precisions/recalls tradeoff");

plt.subplot(4, 2, 3)
fpr, tpr, thresholds = roc_curve(y_test, lr_clf.predict(X_test))
plot_roc_curve(fpr, tpr)

```



```

[45]: scores_dict = {
    'Logistic Regression': {
        'Train': accuracy_score(y_train, lr_clf.predict(X_train)),
        'Test': accuracy_score(y_test, lr_clf.predict(X_test)),
    },

```

```
}
```

Now you can simply select the threshold value that gives you the best precision/recall tradeoff for your task. In our case we need to predict if an employee will absent accurately. so we need to increase

```
[46]: scores = cross_val_score(lr_clf, X_train, y_train, cv=10)

print(f"Scores: {scores}")
print(f"Cross-Validation score mean: {scores.mean() * 100:.2f}% (+/-{scores.
↪std() * 100:.2f})")
```

```
Scores: [0.75555556 0.75555556 0.62222222 0.82222222 0.73333333 0.75555556
0.66666667 0.79545455 0.72727273 0.77272727]
Cross-Validation score mean: 74.07% (+/-5.58)
```

```
[47]: feature_name = X.columns
summary_table = pd.DataFrame(columns=["Features_name"], data=feature_name)
summary_table["Coefficients"] = np.transpose(lr_clf.coef_)
summary_table
```

```
[47]:
```

	Features_name	Coefficients
0	Month_of_absence	0.15
1	Day_of_the_week	-0.20
2	Seasons	-0.22
3	Transportation_expense	0.56
4	Distance_from_Residence_to_Work	-0.11
5	Service_time	-0.03
6	Age	-0.27
7	Work_load_Average/day_	0.10
8	Hit_target	-0.09
9	Disciplinary_failure	-0.98
10	Son	0.63
11	Social_drinker	0.27
12	Social_smoker	-0.01
13	Pet	-0.31
14	Weight	0.25
15	Height	-0.08
16	Body_mass_index	0.00
17	Reason_for_absence_Group_1	0.00
18	Reason_for_absence_Group_2	-0.17
19	Reason_for_absence_Group_3	0.05
20	Reason_for_absence_Group_4	-0.90
21	Education_1	0.00

```
[48]: summary_table.index = summary_table.index + 1
summary_table.loc[0] = ['Intercept', lr_clf.intercept_[0]]
summary_table.sort_index(inplace=True)
```

```
[49]: summary_table["Odds_ratio"] = np.exp(summary_table.Coefficients)
summary_table.sort_values(by="Odds_ratio", ascending=False)
```

```
[49]:
```

	Features_name	Coefficients	Odds_ratio
11	Son	0.63	1.87
4	Transportation_expense	0.56	1.76
12	Social_drinker	0.27	1.31
15	Weight	0.25	1.28
1	Month_of_absence	0.15	1.16
8	Work_load_Average/day_	0.10	1.11
20	Reason_for_absence_Group_3	0.05	1.05
17	Body_mass_index	0.00	1.00
18	Reason_for_absence_Group_1	0.00	1.00
22	Education_1	0.00	1.00
13	Social_smoker	-0.01	0.99
6	Service_time	-0.03	0.97
16	Height	-0.08	0.93
9	Hit_target	-0.09	0.91
5	Distance_from_Residence_to_Work	-0.11	0.90
19	Reason_for_absence_Group_2	-0.17	0.84
2	Day_of_the_week	-0.20	0.82
3	Seasons	-0.22	0.80
7	Age	-0.27	0.77
0	Intercept	-0.27	0.76
14	Pet	-0.31	0.73
21	Reason_for_absence_Group_4	-0.90	0.41
10	Disciplinary_failure	-0.98	0.38

6 Random Forest

```
[50]: print("\n=====RANDOM FOREST=====")
n_estimators = [100, 500, 1000, 1500]
max_features = ['auto', 'sqrt']
max_depth = [2, 5, 10, 15, None]
min_samples_split = [2, 5, 10]
min_samples_leaf = [1, 2, 4]
bootstrap = [True, False]
criterion = ['gini', 'entropy']

params_grid = {
    'n_estimators': n_estimators,
    'max_features': max_features,
    'max_depth': max_depth,
    'min_samples_split': min_samples_split,
    'min_samples_leaf': min_samples_leaf,
    'bootstrap': bootstrap
}
```

```

}

rf_clf = RandomForestClassifier(random_state=42)

rf_cv = GridSearchCV(rf_clf, params_grid, scoring="f1", cv=5, verbose=1,
    ↪n_jobs=-1)

rf_cv.fit(X_train, y_train)
best_params = rf_cv.best_params_
print(f"Best parameters: {best_params}")

rf_clf = RandomForestClassifier(**best_params)
rf_clf.fit(X_train, y_train)

evaluate(rf_clf, X_train, X_test, y_train, y_test)

```

=====RANDOM FOREST=====

Fitting 5 folds for each of 720 candidates, totalling 3600 fits

C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection_validation.py:516: FitFailedWarning: 1800 fits failed out of a total of 3600.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

866 fits failed with the following error:

Traceback (most recent call last):

File "C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection_validation.py", line 859, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py", line 1356, in wrapper
estimator._validate_params()

File "C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py", line 469, in _validate_params
validate_parameter_constraints(

File "C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\utils_param_validation.py", line 98, in
validate_parameter_constraints

raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features'
parameter of RandomForestClassifier must be an int in the range [1, inf), a

float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto' instead.

934 fits failed with the following error:

Traceback (most recent call last):

File "C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection_validation.py", line 859, in _fit_and_score
estimator.fit(X_train, y_train, **fit_params)

File "C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py", line 1356, in wrapper
estimator._validate_params()

File "C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\base.py", line 469, in _validate_params
validate_parameter_constraints(

File "C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\utils_param_validation.py", line 98, in
validate_parameter_constraints

raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto' instead.

warnings.warn(some_fits_failed_message, FitFailedWarning)

C:\Users\maney\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\model_selection_search.py:1135: UserWarning: One or more of the test scores are non-finite: [
nan nan nan nan nan
nan nan

nan	nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan	nan
0.51938376	0.55300069	0.55018444	0.56937874	0.51938376	0.55300069	
0.55018444	0.56937874	0.51938376	0.55115697	0.55018444	0.5654864	
0.51938376	0.55300069	0.55018444	0.5654864	0.51938376	0.55300069	
0.55018444	0.5654864	0.51938376	0.55115697	0.55018444	0.5654864	
0.53108842	0.5582685	0.55394031	0.56937874	0.53108842	0.5582685	
0.55394031	0.56937874	0.53108842	0.56006819	0.55394031	0.56494636	
nan	nan	nan	nan	nan	nan	
nan	nan	nan	nan	nan	nan	
nan	nan	nan	nan	nan	nan	
nan	nan	nan	nan	nan	nan	
nan	nan	nan	nan	nan	nan	
nan	nan	nan	nan	nan	nan	
0.66797106	0.67573025	0.68005013	0.67615779	0.65535434	0.67786466	
0.68131556	0.68145971	0.67628157	0.67445247	0.67961925	0.67756737	

0.65221318	0.6699409	0.66981285	0.66686296	0.66703245	0.67226704
0.65973041	0.66515858	0.66909707	0.67229165	0.67087352	0.66928489
0.66304609	0.66475542	0.66645887	0.66169697	0.66304609	0.66475542
0.66645887	0.66169697	0.67006453	0.66888519	0.66824141	0.66292734
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
0.71099689	0.70631254	0.69796958	0.70252447	0.6957836	0.70831233
0.70201848	0.69700863	0.68952722	0.68766708	0.68854253	0.69313623
0.70235351	0.68543451	0.68520403	0.67690304	0.68691947	0.68017535
0.69233541	0.68611525	0.68849194	0.66255908	0.66760608	0.67663728
0.65742584	0.66522046	0.66346778	0.6750272	0.65742584	0.66522046
0.66346778	0.6750272	0.67089946	0.66222661	0.66201782	0.671172
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
0.68636939	0.70736318	0.70253451	0.69420599	0.70892416	0.70521724
0.70881897	0.70266667	0.69621964	0.70351826	0.69131864	0.69313623
0.6970339	0.69023033	0.69424935	0.68401462	0.69929791	0.68732558
0.68252372	0.68445025	0.69014733	0.67065835	0.67236799	0.67663728
0.65569927	0.65999919	0.66863818	0.67210139	0.65569927	0.65999919
0.66863818	0.67210139	0.66422959	0.66047393	0.66028451	0.67292468
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
0.68599299	0.70736318	0.70108879	0.69789177	0.70892416	0.70154898
0.70690509	0.70542779	0.69458532	0.70170067	0.68965364	0.69313623
0.69890778	0.69023033	0.6926265	0.68401462	0.68924265	0.68732558
0.68252372	0.68634499	0.69014733	0.67065835	0.66926988	0.67663728
0.65569927	0.65999919	0.66522046	0.67210139	0.65569927	0.65999919
0.66522046	0.67210139	0.66422959	0.66047393	0.66028451	0.67292468
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
0.55959633	0.54694873	0.56230565	0.55846879	0.55959633	0.54694873
0.56230565	0.55846879	0.55959633	0.54694873	0.56230565	0.55846879

0.55959633	0.54694873	0.56414937	0.55846879	0.55959633	0.54694873
0.56414937	0.55846879	0.55959633	0.54694873	0.56414937	0.55846879
0.56425903	0.55107516	0.56230565	0.55846879	0.56425903	0.55107516
0.56230565	0.55846879	0.56425903	0.55107516	0.56230565	0.55846879
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
0.68894475	0.69152438	0.69046481	0.68668211	0.69451941	0.68789295
0.68684034	0.68663659	0.68129091	0.69145982	0.69028571	0.68845923
0.68382688	0.68840328	0.68154088	0.68449864	0.68491708	0.68668211
0.6847127	0.68639337	0.67261462	0.68637323	0.68154088	0.68499965
0.67575395	0.67947855	0.68295832	0.68463899	0.67575395	0.67947855
0.68295832	0.68463899	0.67416645	0.67949478	0.67747917	0.67565723
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
0.69008445	0.69833598	0.69833598	0.69142176	0.69662894	0.69317413
0.7010302	0.70088216	0.70454817	0.69913221	0.70441496	0.70259737
0.70266185	0.70747133	0.70582736	0.70218329	0.6857241	0.69686874
0.69951539	0.69765059	0.70027418	0.6984719	0.69853511	0.7037527
0.68333	0.68302182	0.69086498	0.6880811	0.68333	0.68302182
0.69086498	0.6880811	0.69573708	0.68663698	0.6832562	0.67868003
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
0.67964593	0.70353004	0.69445901	0.69441219	0.7128794	0.69129999
0.69611579	0.68783845	0.70614801	0.69921471	0.69739712	0.69739712
0.70760644	0.69606183	0.69765673	0.69928581	0.69196944	0.69226792
0.68878198	0.69202768	0.70573612	0.69497562	0.69322123	0.69497562
0.68341684	0.68249276	0.68249276	0.68597536	0.68341684	0.68249276
0.68249276	0.68597536	0.68883916	0.68330364	0.68608753	0.67868003
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
nan	nan	nan	nan	nan	nan
0.69437282	0.68681577	0.69457064	0.69470734	0.70669422	0.69293045
0.6924437	0.68954732	0.70281461	0.7008866	0.69739712	0.69739712

```

0.70279378 0.69606183 0.69407778 0.69602241 0.69011834 0.69226792
0.68878198 0.6938048 0.70744494 0.69497562 0.69679321 0.69497562
0.68341684 0.68249276 0.68249276 0.68597536 0.68341684 0.68249276
0.68249276 0.68597536 0.68883916 0.68330364 0.68608753 0.67868003]
warnings.warn(

Best parameters: {'bootstrap': False, 'max_depth': 15, 'max_features': 'sqrt',
'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}
TRAINIG RESULTS:
=====
CONFUSION MATRIX:
[[242  6]
 [ 5 194]]
ACCURACY SCORE:
0.9754
CLASSIFICATION REPORT:

```

	0	1	accuracy	macro avg	weighted avg
precision	0.98	0.97	0.98	0.97	0.98
recall	0.98	0.97	0.98	0.98	0.98
f1-score	0.98	0.97	0.98	0.98	0.98
support	248.00	199.00	0.98	447.00	447.00

```

TESTING RESULTS:
=====
CONFUSION MATRIX:
[[87 12]
 [33 60]]
ACCURACY SCORE:
0.7656
CLASSIFICATION REPORT:

```

	0	1	accuracy	macro avg	weighted avg
precision	0.72	0.83	0.77	0.78	0.78
recall	0.88	0.65	0.77	0.76	0.77
f1-score	0.79	0.73	0.77	0.76	0.76
support	99.00	93.00	0.77	192.00	192.00

```

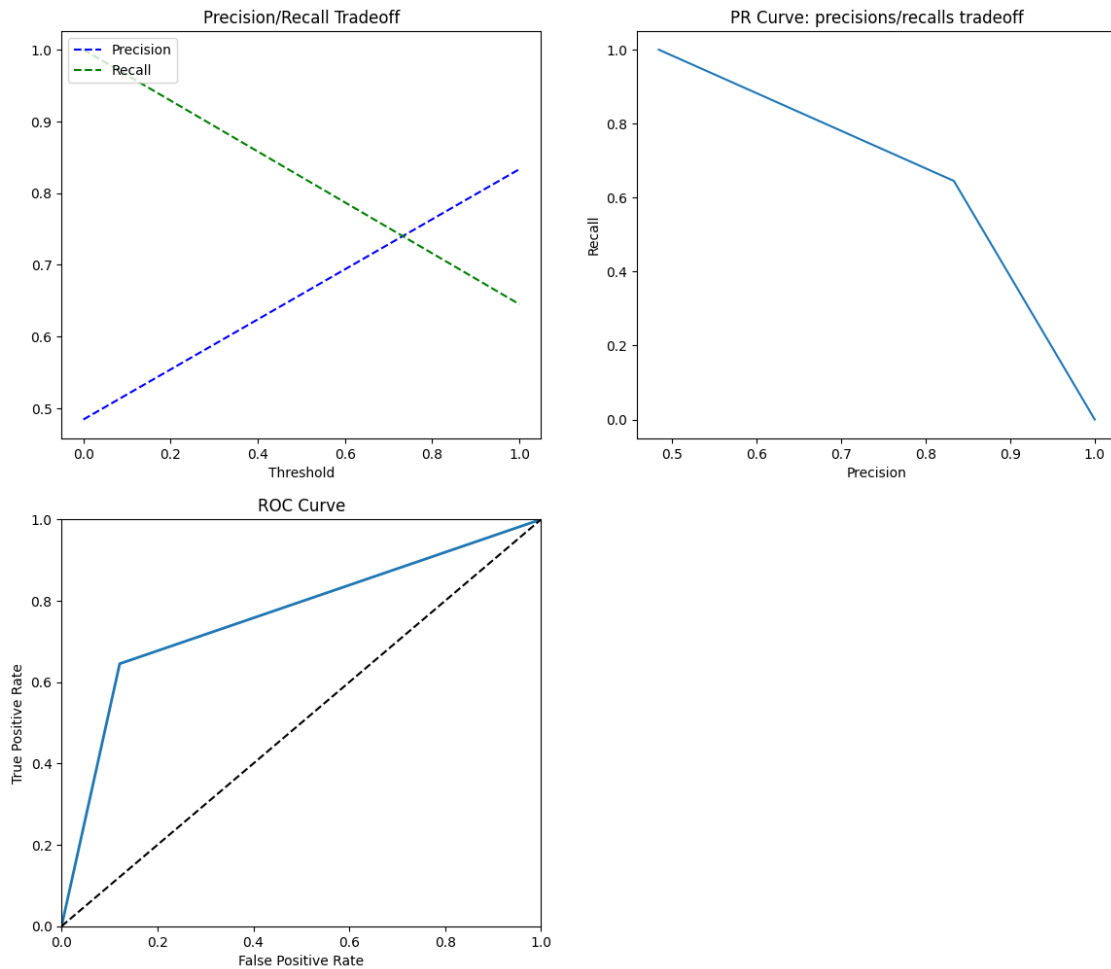
[51]: precisions, recalls, thresholds = precision_recall_curve(y_test, rf_clf.
      ↪predict(X_test))
plt.figure(figsize=(14, 25))
plt.subplot(4, 2, 1)
plot_precision_recall_vs_threshold(precisions, recalls, thresholds)

plt.subplot(4, 2, 2)
plt.plot(precisions, recalls)
plt.xlabel("Precision")
plt.ylabel("Recall")
plt.title("PR Curve: precisions/recalls tradeoff");

plt.subplot(4, 2, 3)

```

```
fpr, tpr, thresholds = roc_curve(y_test, rf_clf.predict(X_test))
plot_roc_curve(fpr, tpr)
```



```
[52]: scores_dict['Random Forest'] = {
        'Train': accuracy_score(y_train, rf_clf.predict(X_train)),
        'Test': accuracy_score(y_test, rf_clf.predict(X_test)),
    }
```

7 XGBoost

```
[53]: n_estimators = [50, 100, 250]
max_depth = [2, 3, 5, 10, 15]
# booster = ['gbtree', 'gblinear']
base_score = [0.2, 0.25, 0.5, 0.75, 0.99]
learning_rate = [0.05, 0.1, 0.5, 0.9, 1]
min_child_weight = [1, 2, 3, 4]
```

```

subsample = [0.5, 0.75, 0.85, 0.9, 1]
colsample_bytree = [0.5, 0.75, 0.85, 0.9, 1]
colsample_bynode = [0.5, 0.75, 0.85, 0.9, 1]
colsample_bylevel = [0.5, 0.75, 0.85, 0.9, 1]

params_grid = {
    'n_estimators': n_estimators,
    'max_depth': max_depth,
    'learning_rate' : learning_rate,
    'min_child_weight' : min_child_weight,
    # 'booster' : booster,
    'base_score' : base_score,
    'subsample': subsample,
    # 'colsample_bytree': colsample_bytree,
    # 'colsample_bynode': colsample_bynode,
    # 'colsample_bylevel': colsample_bylevel,
}

xgb_clf = xgb.XGBClassifier()

xgb_cv = GridSearchCV(xgb_clf, params_grid, cv=5, scoring = 'f1', n_jobs=-1,
    verbose=1)

xgb_cv.fit(X_train, y_train)
best_params = xgb_cv.best_params_
print(f"Best paramters: {best_params}")

xgb_clf = xgb.XGBClassifier(**best_params)
xgb_clf.fit(X_train, y_train)

evaluate(xgb_clf, X_train, X_test, y_train, y_test)

```

Fitting 5 folds for each of 7500 candidates, totalling 37500 fits
 Best paramters: {'base_score': 0.2, 'learning_rate': 0.05, 'max_depth': 5,
 'min_child_weight': 1, 'n_estimators': 250, 'subsample': 1}
 TRAINIG RESULTS:

=====

CONFUSION MATRIX:

```
[[243  5]
 [ 9 190]]
```

ACCURACY SCORE:

0.9687

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.96	0.97	0.97	0.97	0.97
recall	0.98	0.95	0.97	0.97	0.97

f1-score	0.97	0.96	0.97	0.97	0.97
support	248.00	199.00	0.97	447.00	447.00

TESTING RESULTS:

=====

CONFUSION MATRIX:

```
[[83 16]
 [32 61]]
```

ACCURACY SCORE:

0.7500

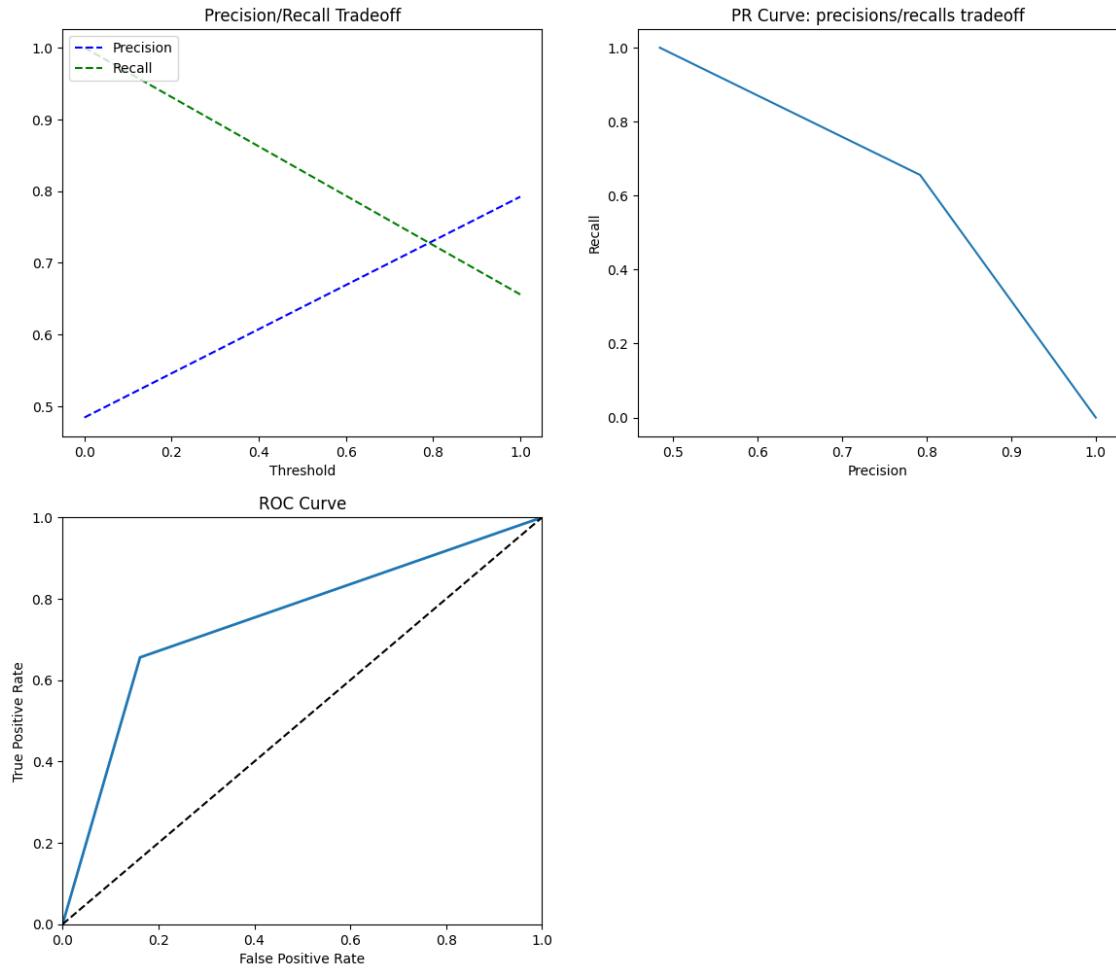
CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.72	0.79	0.75	0.76	0.76
recall	0.84	0.66	0.75	0.75	0.75
f1-score	0.78	0.72	0.75	0.75	0.75
support	99.00	93.00	0.75	192.00	192.00

```
[54]: precisions, recalls, thresholds = precision_recall_curve(y_test, xgb_clf.
      ↪predict(X_test))
plt.figure(figsize=(14, 25))
plt.subplot(4, 2, 1)
plot_precision_recall_vs_threshold(precisions, recalls, thresholds)

plt.subplot(4, 2, 2)
plt.plot(precisions, recalls)
plt.xlabel("Precision")
plt.ylabel("Recall")
plt.title("PR Curve: precisions/recalls tradeoff");

plt.subplot(4, 2, 3)
fpr, tpr, thresholds = roc_curve(y_test, xgb_clf.predict(X_test))
plot_roc_curve(fpr, tpr)
```



```
[55]: scores_dict['XGBoost'] = {
        'Train': accuracy_score(y_train, xgb_clf.predict(X_train)),
        'Test': accuracy_score(y_test, xgb_clf.predict(X_test)),
    }
```

8 Support Vector Machine

```
[56]: param_grid = {
        'C': [60, 70, 75, 65 ],
        'gamma': [0.002, 0.001, 0.0009, 0.0008, 0.0007],
        'kernel': ['rbf', 'poly', 'linear']
    }

svm_cv = GridSearchCV(SVC(), param_grid, scoring='f1', verbose=1, cv=5)
svm_cv.fit(X_train, y_train)
```

```

best_params = svm_cv.best_params_
print(f"Best params: {best_params}")

svm_clf = SVC(**best_params)
svm_clf.fit(X_train, y_train)
evaluate(svm_clf, X_train, X_test, y_train, y_test)

```

Fitting 5 folds for each of 60 candidates, totalling 300 fits

Best params: {'C': 60, 'gamma': 0.002, 'kernel': 'linear'}

TRAINING RESULTS:

=====

CONFUSION MATRIX:

```

[[193  55]
 [ 56 143]]

```

ACCURACY SCORE:

0.7517

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.78	0.72	0.75	0.75	0.75
recall	0.78	0.72	0.75	0.75	0.75
f1-score	0.78	0.72	0.75	0.75	0.75
support	248.00	199.00	0.75	447.00	447.00

TESTING RESULTS:

=====

CONFUSION MATRIX:

```

[[84 15]
 [33 60]]

```

ACCURACY SCORE:

0.7500

CLASSIFICATION REPORT:

	0	1	accuracy	macro avg	weighted avg
precision	0.72	0.80	0.75	0.76	0.76
recall	0.85	0.65	0.75	0.75	0.75
f1-score	0.78	0.71	0.75	0.75	0.75
support	99.00	93.00	0.75	192.00	192.00

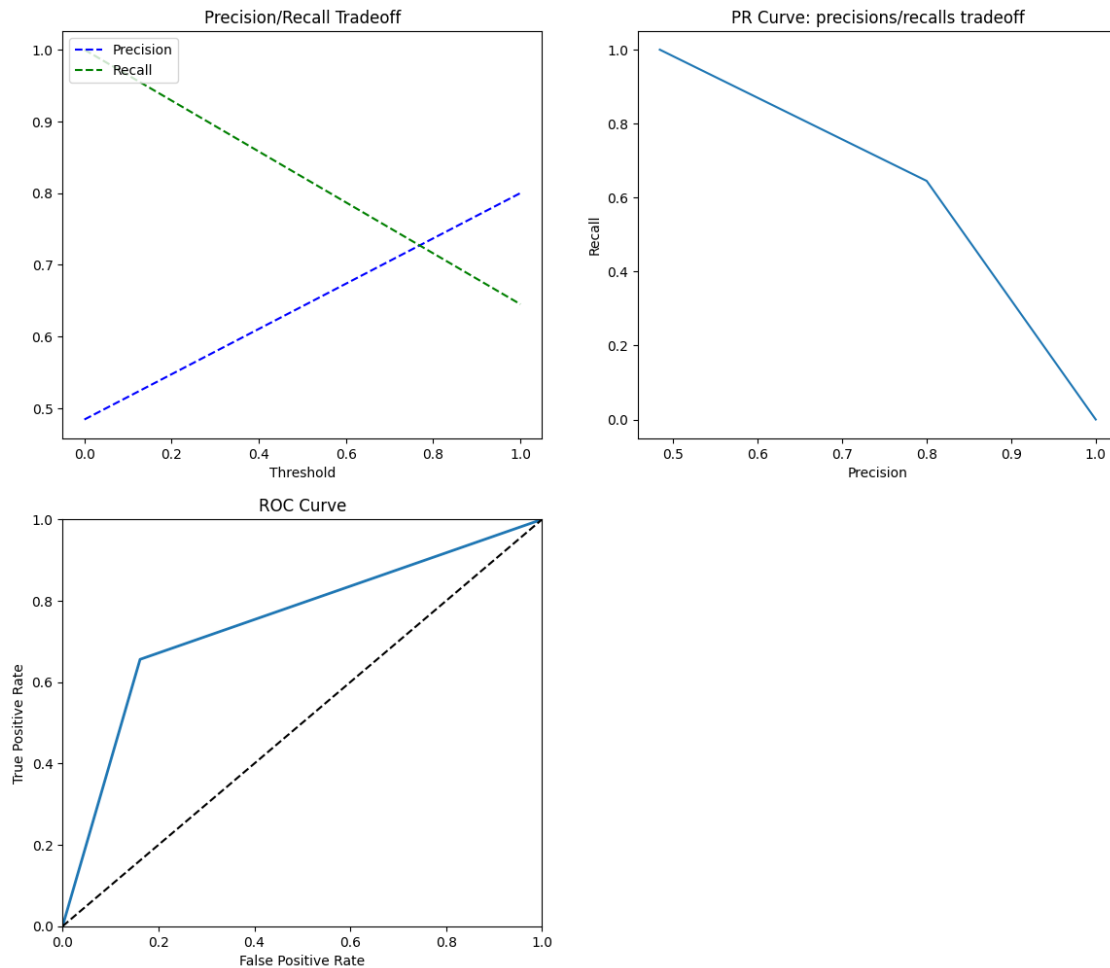
```

[57]: precisions, recalls, thresholds = precision_recall_curve(y_test, svm_clf.
      ↪predict(X_test))
plt.figure(figsize=(14, 25))
plt.subplot(4, 2, 1)
plot_precision_recall_vs_threshold(precisions, recalls, thresholds)

plt.subplot(4, 2, 2)
plt.plot(precisions, recalls)
plt.xlabel("Precision")
plt.ylabel("Recall")
plt.title("PR Curve: precisions/recalls tradeoff");

```

```
plt.subplot(4, 2, 3)
fpr, tpr, thresholds = roc_curve(y_test, xgb_clf.predict(X_test))
plot_roc_curve(fpr, tpr)
```



```
[58]: scores_dict['Support Vector Machine'] = {
        'Train': accuracy_score(y_train, svm_clf.predict(X_train)),
        'Test': accuracy_score(y_test, svm_clf.predict(X_test)),
    }
```

9 Comparing Machine Learning models

Area Under the Curve score (AUC) is good way to compare classifiers. A perfect classifier AUC will have a ROC AUC equal to 1.



```
[59]: ml_models = {
        'Logistic Regression': lr_clf,
        'Random Forest': rf_clf,
        'XGboost': xgb_clf,
        'Support Vector Machine': svm_clf
    }
    for model in ml_models:
        print(f"{model.upper()} roc_auc_score: {roc_auc_score(y_test,
        ↪ml_models[model].predict(X_test)):.3f}")
```

```
LOGISTIC REGRESSION roc_auc_score: 0.731
RANDOM FOREST roc_auc_score: 0.762
XGBOOST roc_auc_score: 0.747
SUPPORT VECTOR MACHINE roc_auc_score: 0.747
```

```
[60]: scores_df = pd.DataFrame(scores_dict)

scores_df.plot(kind='barh', figsize=(15, 8))
```

[60]: <Axes: >



10 Save the models

```
[61]: with open('Xgb_clf', 'wb') as file:
        pickle.dump(xgb_clf, file)

    with open('Rf_clf', 'wb') as file:
```

```
pickle.dump(rf_clf, file)
```